

✦ Member-only story

UNet — Line by Line Explanation

Example UNet Implementation



Jeremy Zhang · Follow

Published in Towards Data Science

4 min read · Oct 18, 2019



Listen



Share



More

UNet, evolved from the traditional convolutional neural network, was first designed and applied in 2015 to process biomedical images. As a general convolutional neural network focuses its task on image classification, where input is an image and output is one label, but in biomedical cases, it requires us not only to distinguish whether there is a disease, but also to localise the area of abnormality.

UNet is dedicated to solving this problem. **The reason it is able to localise and distinguish borders is by doing classification on every pixel, so the input and output share the same size.** For example, for an input image of size 2x2:

```
[[255, 230], [128, 12]] # each number is a pixel
```

the output will have the same size of 2x2:

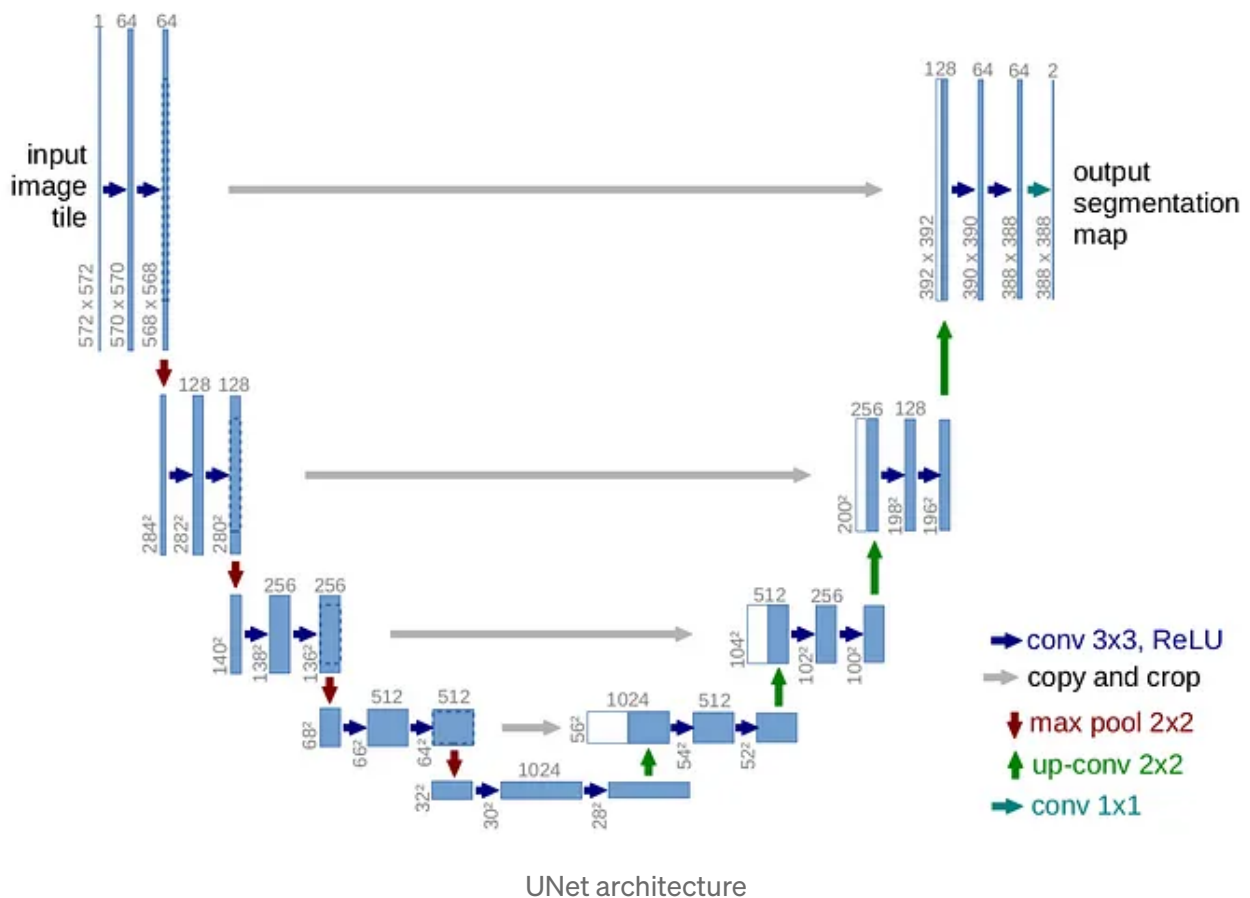
```
[[1, 0], [1, 1]] # could be any number between [0, 1]
```

Now let's get to the detail implementation of UNet. I will:

1. Show the overview of UNet
2. Breakdown the implementation line by line and further explain it

Overview

The network has basic foundation looks like:



First sight, it has a “U” shape. The architecture is symmetric and consists of two major parts — the left part is called contracting path, which is constituted by the general convolutional process; the right part is expansive path, which is constituted by transposed 2d convolutional layers (you can think it as an upsampling technic for now).

Now let's have a quick look at the implementation:

```

1  def build_model(input_layer, start_neurons):
2      conv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(input_layer)
3      conv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(conv1)
4      pool1 = MaxPooling2D((2, 2))(conv1)
5      pool1 = Dropout(0.25)(pool1)
6
7      conv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(pool1)
8      conv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(conv2)
9      pool2 = MaxPooling2D((2, 2))(conv2)
10     pool2 = Dropout(0.5)(pool2)
11
12     conv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(pool2)
13     conv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(conv3)
14     pool3 = MaxPooling2D((2, 2))(conv3)
15     pool3 = Dropout(0.5)(pool3)
16
17     conv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(pool3)
18     conv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(conv4)
19     pool4 = MaxPooling2D((2, 2))(conv4)
20     pool4 = Dropout(0.5)(pool4)
21
22     # Middle
23     convm = Conv2D(start_neurons * 16, (3, 3), activation="relu", padding="same")(pool4)
24     convm = Conv2D(start_neurons * 16, (3, 3), activation="relu", padding="same")(convm)
25
26     deconv4 = Conv2DTranspose(start_neurons * 8, (3, 3), strides=(2, 2), padding="same")(convm)
27     uconv4 = concatenate([deconv4, conv4])
28     uconv4 = Dropout(0.5)(uconv4)
29     uconv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(uconv4)
30     uconv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(uconv4)
31
32     deconv3 = Conv2DTranspose(start_neurons * 4, (3, 3), strides=(2, 2), padding="same")(uconv4)
33     uconv3 = concatenate([deconv3, conv3])
34     uconv3 = Dropout(0.5)(uconv3)
35     uconv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(uconv3)
36     uconv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(uconv3)
37
38     deconv2 = Conv2DTranspose(start_neurons * 2, (3, 3), strides=(2, 2), padding="same")(uconv3)
39     uconv2 = concatenate([deconv2, conv2])
40     uconv2 = Dropout(0.5)(uconv2)
41     uconv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(uconv2)
42     uconv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(uconv2)
43
44     deconv1 = Conv2DTranspose(start_neurons * 1, (3, 3), strides=(2, 2), padding="same")(uconv2)
45     uconv1 = concatenate([deconv1, conv1])
46     uconv1 = Dropout(0.5)(uconv1)
47     uconv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(uconv1)
48     uconv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(uconv1)

```

```

48     uconv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same", uconv1)
49
50     output_layer = Conv2D(1, (1,1), padding="same", activation="sigmoid")(uconv1)
51
52     return output_layer
53
54     input_layer = Input((img_size_target, img_size_target, 1))
55     output_layer = build_model(input_layer, 16)

```

unet.py hosted with ❤ by GitHub

[view raw](#)

The code is referred from a [kernel](#) of Kaggle competition, in general, most UNet follows the same structure.

Now let's break down the implementation line by line and maps to the corresponding parts on the image of UNet architecture.

Line by Line Explanation

Contracting Path

The contracting path follows the formula:

conv_layer1 -> conv_layer2 -> max_pooling -> dropout(optional)

So the first part of our code is:

```


1 conv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(input_layer)
2 conv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(conv1)
3 pool1 = MaxPooling2D((2, 2))(conv1)
4 pool1 = Dropout(0.25)(pool1)

```

unet2.py hosted with ❤ by GitHub


[view raw](#)

which matches to:



Notice that each process constitutes two convolutional layers, and the number of channel changes from $1 \rightarrow 64$, as convolution process will increase the depth of the image. The red arrow pointing down is the max pooling process which halves down size of image (the size reduced from $572 \times 572 \rightarrow 568 \times 568$ is due to padding issues, but the implementation here uses padding= “same”).

The process is repeated 3 more times:



with code:

```

1 conv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(pool1)
2 conv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(conv2)
3 pool2 = MaxPooling2D((2, 2))(conv2)
4 pool2 = Dropout(0.5)(pool2)
5
6 conv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(pool2)
7 conv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(conv3)
8 pool3 = MaxPooling2D((2, 2))(conv3)
9 pool3 = Dropout(0.5)(pool3)
10
11 conv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(pool3)
12 conv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(conv4)
13 pool4 = MaxPooling2D((2, 2))(conv4)
14 pool4 = Dropout(0.5)(pool4)

```

UNET3.py hosted with ❤ by GitHub

[view raw](#)

and now we reaches at the bottommost:



still 2 convolutional layers are built, but with no max pooling:

```
1 # Middle
2 convm = Conv2D(start_neurons * 16, (3, 3), activation="relu", padding="same")(pool4)
3 convm = Conv2D(start_neurons * 16, (3, 3), activation="relu", padding="same")(convm)
```

UNET4.py hosted with ❤ by GitHub


[view raw](#)

The image at this moment has been resized to 28x28x1024. Now let's get to the expansive path.

Expansive Path

In the expansive path, the image is going to be upsized to its original size. The formula follows:

conv_2d_transpose -> concatenate -> conv_layer1 -> conv_layer2



```
1 deconv4 = Conv2DTranspose(start_neurons * 8, (3, 3), strides=(2, 2), padding="same")(conv4)
2 uconv4 = concatenate([deconv4, conv4])
3 uconv4 = Dropout(0.5)(uconv4)
4 uconv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(uconv4)
5 uconv4 = Conv2D(start_neurons * 8, (3, 3), activation="relu", padding="same")(uconv4)
```

UNET5.py hosted with ❤ by GitHub

[view raw](#)

Transposed convolution is an upsampling technic that expands the size of images. There is a visualised demo [here](#) and an explanation [here](#). Basically, it does some padding on the original image followed by a convolution operation.

After the transposed convolution, the image is upsized from 28x28x1024 → 56x56x512, and then, this image is concatenated with the corresponding image from the contracting path and together makes an image of size 56x56x1024. The reason

here is to combine the information from the previous layers in order to get a more precise prediction.

In line 4 and line 5, 2 other convolution layers are added.

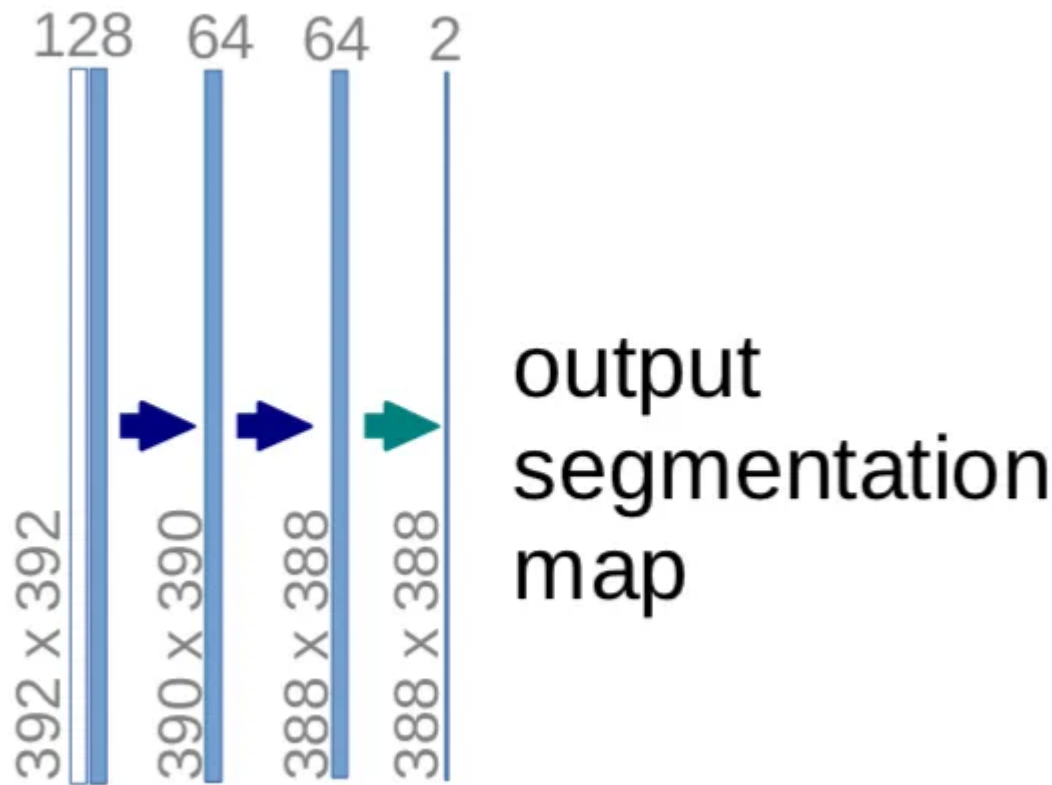
Same as before, this process is repeated 3 more times:

```
1 deconv3 = Conv2DTranspose(start_neurons * 4, (3, 3), strides=(2, 2), padding="same")(u
2 uconv3 = concatenate([deconv3, conv3])
3 uconv3 = Dropout(0.5)(uconv3)
4 uconv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(uconv3)
5 uconv3 = Conv2D(start_neurons * 4, (3, 3), activation="relu", padding="same")(uconv3)
6
7 deconv2 = Conv2DTranspose(start_neurons * 2, (3, 3), strides=(2, 2), padding="same")(u
8 uconv2 = concatenate([deconv2, conv2])
9 uconv2 = Dropout(0.5)(uconv2)
10 uconv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(uconv2)
11 uconv2 = Conv2D(start_neurons * 2, (3, 3), activation="relu", padding="same")(uconv2)
12
13 deconv1 = Conv2DTranspose(start_neurons * 1, (3, 3), strides=(2, 2), padding="same")(u
14 uconv1 = concatenate([deconv1, conv1])
15 uconv1 = Dropout(0.5)(uconv1)
16 uconv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(uconv1)
17 uconv1 = Conv2D(start_neurons * 1, (3, 3), activation="relu", padding="same")(uconv1)
```

unet6.py hosted with ❤ by GitHub

[view raw](#)

Now we've reached the uppermost of the architecture, the last step is to reshape the image to satisfy our prediction requirements.



```
1 output_layer = Conv2D(1, (1,1), padding="same", activation="sigmoid")(uconv1)
```

unet7.py hosted with ❤ by GitHub

[view raw](#)

The last layer is a convolution layer with 1 filter of size 1x1 (notice that there is no dense layer in the whole network). And the rest left is the same for neural network training.

Conclusion

UNet is able to do image localisation by predicting the image pixel by pixel and the author of UNet claims in his [paper](#) that the network is strong enough to do good prediction based on even few data sets by using excessive data augmentation techniques. There are many applications of image segmentation using UNet and it also occurs in lots of competitions. One should try out on yourself and I hope this post could be a good starting point for you.

Reference:

- <https://github.com/hlamba28/UNET-TGS/blob/master/TGS%20UNET.ipynb>
- <https://towardsdatascience.com/understanding-semantic-segmentation-with-unet-6be4f42d4b47>

- <https://towardsdatascience.com/types-of-convolutions-in-deep-learning-717013397f4d>
- <https://medium.com/activating-robotic-minds/up-sampling-with-transposed-convolution-9ae4f2df52d0>
- <https://www.kaggle.com/phoenigs/u-net-dropout-augmentation-stratification>

Machine Learning

Neural Networks

Data Science



Follow

Written by Jeremy Zhang

1.2K Followers · Writer for Towards Data Science

Hmm...I am a data scientist looking to catch up the tide...

More from Jeremy Zhang and Towards Data Science



Jeremy Zhang in Towards Data Science

Reinforcement Learning—Implement TicTacToe

We have implemented grid world game by iteratively updating Q value function, which is the estimating value of (state, action) pair. This...

6 min read · May 19, 2019



271



8



Open in app ↗



Search



Thu Vu in Towards Data Science

How to Learn AI on Your Own (a self-study guide)

If your hands touch a keyboard for work, Artificial Intelligence is going to change your job in the next few years.

★ · 12 min read · Jan 5

👏 2.2K 💬 24

🔖⁺ ⋮



 Michael Berk in Towards Data Science

1.5 Years of Spark Knowledge in 8 Tips

My learnings from Databricks customer engagements

8 min read · Dec 24, 2023

👏 1.4K 💬 8

🔖⁺ ⋮



 Jeremy Zhang in Towards Data Science

Importance Sampling Introduction

Estimate Expectations from a Different Distribution

4 min read · Aug 31, 2019



432



1



See all from Jeremy Zhang

See all from Towards Data Science

Recommended from Medium




FIGURE 4 The overview of the proposed U-Net based B2U2S architecture for lung segmentation

Alexquesada


U-Net: A Versatile Deep Learning Architecture for Image Segmentation

This article is one of the assignments of Data Glacier Intership 2023.

6 min read · Aug 4, 2023

18

...



Roushanak Rahmat, PhD in Code Like A Girl

U-Net vs Residual U-Net vs Attention U-Net vs Attention Residual U-Net


UNet

🌟 · 6 min read · Aug 7, 2023

👏 50 💬

🔖⁺ ⋮

Lists



Predictive Modeling w/ Python

20 stories · 824 saves



Practical Guides to Machine Learning

10 stories · 953 saves




Natural Language Processing


1117 stories · 589 saves



data science and AI

39 stories · 49 saves



 Shameerayaseen

U-Net: Advancing Image Segmentation with Convolutional Neural Networks

What is Image segmentation?

4 min read · Sep 4, 2023



61

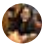
 Samruddhi Chitnis

Image Segmentation Transfer Learning Part 3

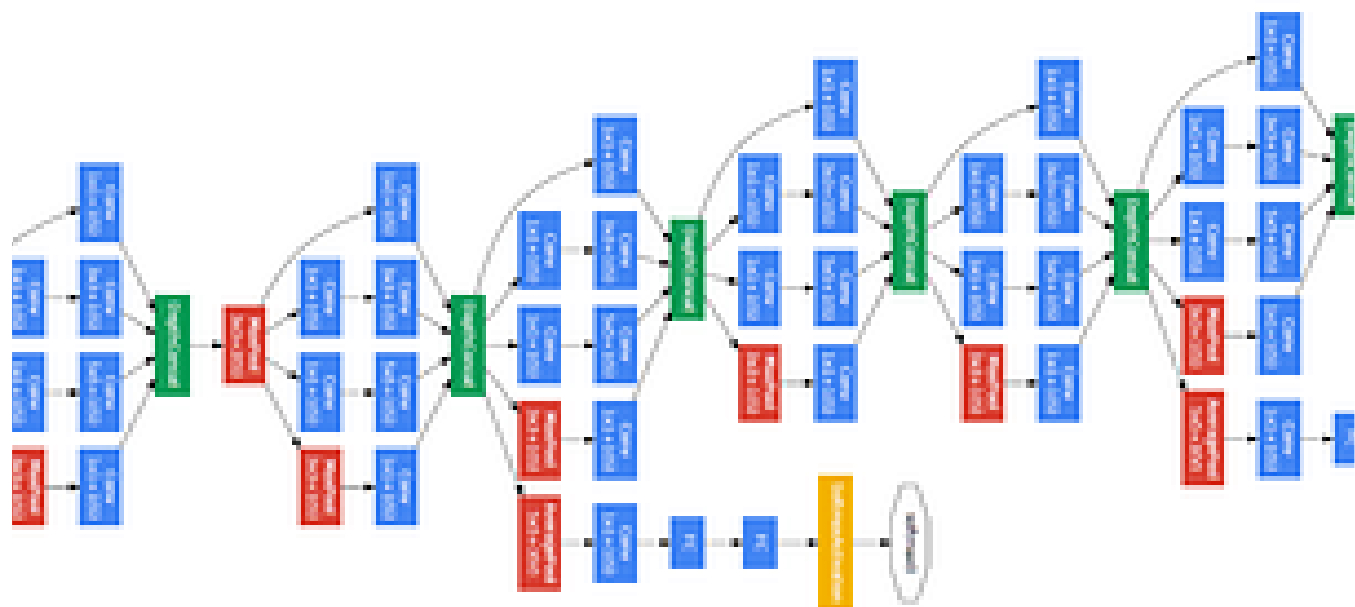
Table of Contents:

3 min read · Oct 11, 2023



42





 Everton Gomedede, PhD

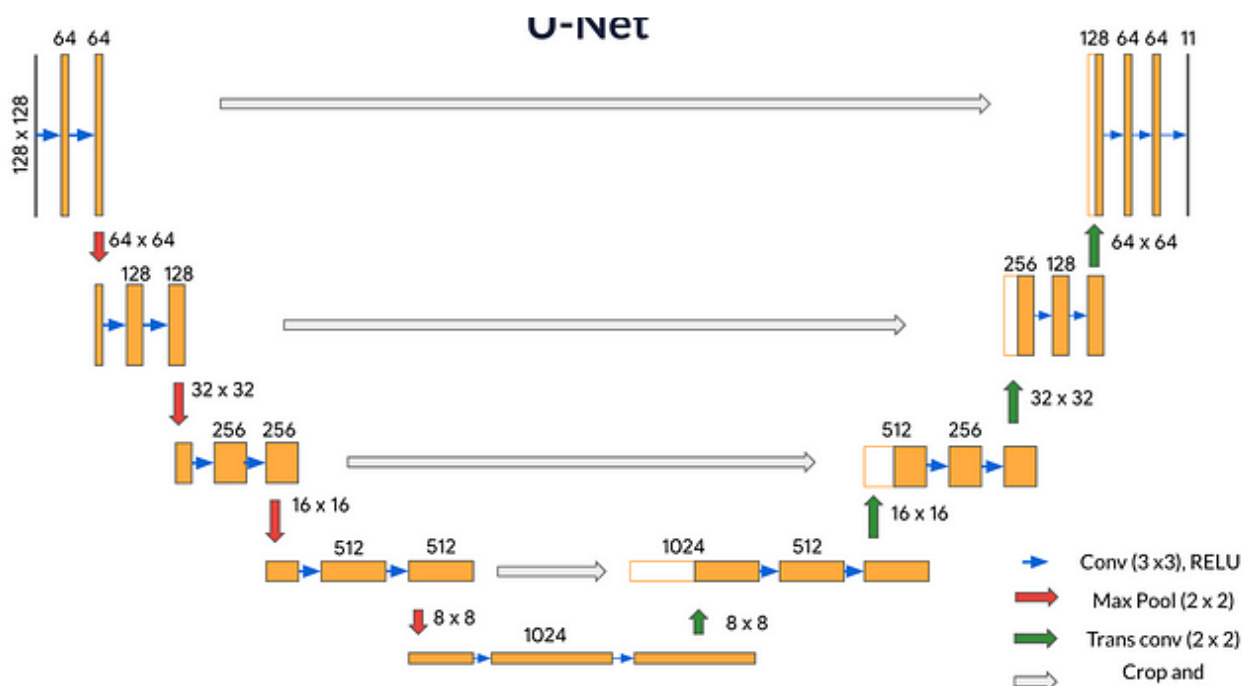
Exploring GoogLeNet: A Revolutionary Deep Learning Architecture

Introduction

7 min read · Oct 2, 2023

 63 



 Lukman Aliyu

U-Net Architecture: Revolutionizing Computer Vision Through Innovative Image Segmentation

Introduction

5 min read · Aug 28, 2023



See more recommendations