

Differentiable Sampling and Argmax

⋮

WIP, last updated: 2019.12.6

Introduction

Softmax is a commonly used function for turning an **unnormalized log probability** into a normalized probability (or **categorical distribution**).

$$\pi = \text{softmax}(\mathbf{o}) = \frac{e^{o_j}}{\sum_j e^{o_j}},$$
$$o_j \in (-\infty, +\infty)$$

Say \mathbf{o} is the output of a neural network before softmax, we call \mathbf{o} the **unnormalized log probability**.

After softmax, we usually **sample** from this categorical distribution, or taking an **argmax** function to select the index. However, one can notice that neither the **sampling** nor the **argmax** is **differentiable**.

Researchers have proposed several works to make this possible. I am going to discuss them here.

Sampling

I will introduce Gumbel Softmax [\[1611.01144\]](#), which have made the **sampling** procedure differentiable.

Gumbel Max

First, we need to introduce **Gumbel Max**. In short, Gumbel Max is a trick to use gumbel distribution to sample a categorical distribution.

Say we want to sample from a categorical distribution π .

The usual way of doing this is using π to separate $[0, 1]$ into intervals, sampling from a uniform distribution $U \sim [0, 1]$, and see where it locates.

The Gumbel Max trick provides an alternative way of doing this. It use **Reparameterization Trick** to avoid the stochastic node during backpropagation.

$$y = \arg \max_i (o_i + g_i)$$

where $g_i \sim \text{Gumbel}(0, 1)$, which can be sampled by $-\log(-\log(\text{Uniform}[0, 1]))$. We can prove that y is distributed according to π .

Prove that

$y = \arg \max_i (o_i + g_i)$, where $g_i \sim \text{Gumbel}(0, 1)$
which can be sampled by $-\log(-\log(\text{Uniform}[0, 1]))$
is distributed with $\pi = \text{softmax}(o_i) = \frac{e^{o_i}}{\sum_j e^{o_j}}$

Prerequisites

Gumbel Distribution (param by location **** μ , and scale $\beta > 0$) ([wikipedia](#))

CDF: $F(x; \mu, \beta) = e^{-e^{(x-\mu)/\beta}}$

PDF: $f(x; \mu, \beta) = \frac{1}{\beta} e^{-(z+e^{-z})}, z = \frac{x-\mu}{\beta}$

Mean: $E(X) = \mu + \gamma\beta$, $\gamma \approx 0.5772$ is the Euler–Mascheroni constant.

Quantile Function: $Q(p) = \mu - \beta \log(-\log(p))$ (Quantile Function is used to sample random variables from a distribution given CDF, it is also called inverse CDF)

Proof

We actually want to prove that $\text{Gumbel}(\mu = o_i, \beta = 1)$

is distributed with $\pi_i = \frac{e^{o_i}}{\sum_j e^{o_j}}$.

We can find that $\text{Gumbel}(\mu = o_i, \beta = 1)$ has the following PDF and CDF

$$\begin{aligned} f(x; \mu, 1) &= e^{-(x-\mu)-e^{-(x-\mu)}} \\ F(x; \mu, 1) &= e^{-e^{-(x-\mu)}} \end{aligned} \quad (2)$$

.Then, the probability that all other $\pi_{j \neq i}$ are less than π_i is:

$$\Pr(\pi_i \text{ is the largest} | \pi_i, \{o_j\}) = \prod_{j \neq i} e^{-e^{-(\pi_i - o_j)}}$$

We know the marginal distribution over π_i and we are able to integrate it out to find the overall probability: $p(x) = \int_y p(x, y) dy = \int_y p(x|y)p(y) dy$

$$\Pr(i \text{ is largest} | \{o_j\}) = \int e^{-(\pi_i - o_i) - e^{-(\pi_i - o_i)}} \times \prod_{j \neq i} e^{-e^{-(\pi_i - o_j)}} d\pi_i$$

$$= \int e^{-\pi_i + o_i - e^{-\pi_i} \sum_j e^{o_j}} d\pi_i \quad (4)$$

$$= \frac{e^{o_i}}{\sum_j e^{o_j}} \quad (5)$$

which is exactly a softmax probability. QED.

Reference: <https://lips.cs.princeton.edu/the-gumbel-max-trick-for-discrete-distributions/>****

Gumbel Softmax

Notice that there is still an argmax in Gumbel Max, which still makes it indifferentiable. Therefore, we use a softmax function to approximate this argmax procedure.

$$\mathbf{y} = \frac{e^{(o_i + g_i)/\tau}}{\sum_j e^{(o_j + g_j)/\tau}}$$

where $\tau \in (0, \infty)$ is a temperature hyperparameter.

We note that the output of Gumbel Softmax function here is a vector which sum to 1, which somewhat looks like a one-hot vector (but it's not).

So by far, this does not actually replace the argmax function.

To actually get a pure one-hot vector, we need to use a **Straight-Through (ST) Gumbel Trick**.

Let's directly see an [implementation of Gumbel Softmax in PyTorch](#)

(We use the hard mode, soft mode does not get a pure one-hot vector).

```
def gumbel_softmax(logits, tau=1, hard=False, eps=1e-10, dim=-1):
    # type: (Tensor, float, bool, float, int) -> Tensor
    r"""
    Samples from the Gumbel-Softmax distribution (`Link 1`_ `Link 2`_) a
```

Args:

logits: `[..., num_features]` unnormalized log probabilities
tau: non-negative scalar temperature
hard: if ``True``, the returned samples will be discretized as one-hot but will be differentiated as if it is the soft sample in autograd
dim (int): A dimension along which softmax will be computed. Default

Returns:

Sampled tensor of same shape as `logits` from the Gumbel-Softmax distribution
If ``hard=True``, the returned samples will be one-hot, otherwise they will be probability distributions that sum to 1 across `dim`.

.. note::

This function is here for legacy reasons, may be removed from nn.Module in the future.

.. note::

The main trick for `hard` is to do `y_hard - y_soft.detach() + y_soft`

It achieves two things:

- makes the output value exactly one-hot (since we add then subtract `y_soft` value)
- makes the gradient equal to `y_soft` gradient (since we strip all other gradients)

Examples::

```
>>> logits = torch.randn(20, 32)
>>> # Sample soft categorical using reparametrization trick:
>>> F.gumbel_softmax(logits, tau=1, hard=False)
>>> # Sample hard categorical using "Straight-through" trick:
>>> F.gumbel_softmax(logits, tau=1, hard=True)
```

```

.. _Link 1:
    https://arxiv.org/abs/1611.00712
.. _Link 2:
    https://arxiv.org/abs/1611.01144
"""

if eps != 1e-10:
    warnings.warn("`eps` parameter is deprecated and has no effect.")

```

```

gumbels = -torch.empty_like(logits).exponential_().log() # ~Gumbel(0
gumbels = (logits + gumbels) / tau
When forwarding the code uses an argmax to get an actual one-hot vector.
And it uses ret = y_hard - y_soft.detach() + y_soft, y_hard has no grad,
and by minusing y_soft.detach() and adding y_soft, it achieves a grad from
y_soft without modifying the forwarding value.
# Straight through.

```

```

So eventually, we are able to get a pure one-hot vector in forward pass, and a grad when
back propagating, which makes the sampling procedure differentiable.
index = y_soft.max(dim=-1, keepdim=True)[1]
y_hard = torch.zeros_like(logits).scatter_(dim=-1, index, 1.0)
ret = y_hard - y_soft.detach() + y_soft

```

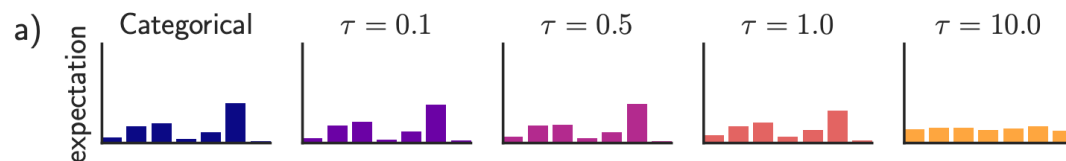
Finally, let's look at how τ affects the sampling procedure. The below image shows the sampling distribution (which is also called the Concrete Distribution [1611.00712]) and one random sample instance when using different hyperparameter τ .

```

# Reparametrization trick
ret = y_soft
return ret

```





b) when $\tau \rightarrow 0$, the softmax becomes an argmax and the Gumbel-Softmax distribution becomes the categorical distribution. During training, we let $\tau > 0$ to allow gradients past the sample, then gradually anneal the temperature τ (but not completely to 0, as the gradients would blow up).

Figure 1: The Gumbel-Softmax distribution interpolates between discrete one-hot-encoded categorical distributions and continuous categorical densities. (a) For low temperatures ($\tau = 0.1, \tau = 0.5$), the expected value of a Gumbel-Softmax random variable approaches the expected value of a categorical random variable with the same logits. As the temperature increases ($\tau = 1.0, \tau = 10.0$), the expected value converges to a uniform distribution over the categories. (b) Samples from Gumbel-Softmax distributions are identical to samples from a categorical distribution as $\tau \rightarrow 0$. At higher temperatures, Gumbel-Softmax samples are no longer one-hot, and become uniform as $\tau \rightarrow \infty$.

How to make argmax differentiable?

image from <https://arxiv.org/abs/1611.01144>

Intuitively, the **Straight-Through Trick** is also applicable for softmax+argmax

[gitbook](#)

[lzhbrian.me](#)

Search

⌘ K

Tzu-Heng's wiki

Tzu-Heng's wiki

MACHINE LEARNING

Traditionals

Deep Learning

Image Classification (CNN)

Detection

Semantic Segmentation

Generative Adversarial Networks

Some have introduced the soft-argmax function. It doesn't actually makes it differentiable, but use a continuous function to approximate the softmax+argmax procedure.

$$\pi = \text{soft-argmax}(\mathbf{o}) = \frac{e^{\beta \mathbf{o}}}{\sum_j e^{\beta o_j}}$$

where β can be a large value to make π very much "look like" a one-hot vector.

Style Transfer

Recommender Systems

Meta Learning

NOTES

Differentiable Sampling and Argmax

GAN theory

Multi-task Learning (MTL)

Disentanglement in GANs



Powered By **GitBook**

Discussion

- Goal
 - **softmax + argmax** is used for classification, we only want the index with the highest probability.
 - **gumbel softmax + argmax** is used for sampling, we may want to sample an index not with the highest probability.
- Deterministic
 - **softmax + argmax** is deterministic. Get the index with the highest probability.
 - **gumbel softmax + argmax** is stochastic. We need to sample from a gumbel distribution in the beginning.
- Output vector
 - **softmax** and **gumbel softmax** both output a vector sum to 1.
 - **softmax** outputs a *normalized probability distribution*.
 - **gumbel softmax** outputs a *sample* somewhat more similar to a one-hot vector. (can be controlled by τ)
- **Straight-Through Trick** can actually be applied to both **softmax + argmax** and **gumbel softmax + argmax**, which can make both of them differentiable. (?)

Reference

- Gumbel Softmax [1611.01144]
- Concrete Distribution (Gumbel Softmax Distribution) [1611.00712]
- Eric Jang official blog: <https://blog.evjang.com/2016/11/tutorial-categorical-variational.html>
- PyTorch Implementation of Gumbel Softmax:
https://pytorch.org/docs/stable/nn.functional.html#torch.nn.functional.gumbel_softmax

- <https://timvieira.github.io/blog/post/2014/07/31/gumbel-max-trick/>
- <https://lips.cs.princeton.edu/the-gumbel-max-trick-for-discrete-distributions/>



Machine Learning - Previous

Meta Learning

Next - Notes

GAN theory



Last modified 1yr ago

WAS THIS PAGE HELPFUL? ☹️ = 😊